

Atharva Chundurwar

Tempe, AZ, USA | atharvachundurwar841@gmail.com | +1(480)743-6844 | [LinkedIn](#) | [GitHub](#)

TECHNICAL SKILLS

Programming: Python, C, C++, SQL, R, PostgreSQL, Shell/Bash, Linux.

Tools & Technologies: LangChain, LangGraph, Hugging Face, Pinecone, FAISS, Chroma DB, Weaviate, Llama-Index, Ollama, AWS, GCP, Azure, Kubernetes, Docker, Git/GitHub, Hadoop, Hive, dbt, Prometheus, Grafana, Flask, FastAPI, Apache Spark, Airflow, Snowflake, MLflow, DVC, Databricks, Tableau, Power BI, Mathematical Statistics and Analytics.

Skills: Machine / Deep Learning, Generative AI, NLP, MLOps, RAG Systems, Data Mining, API Development, LLM Fine-Tuning.

PROFESSIONAL EXPERIENCE

Revmo AI — R&D Intern

Oct 2025 - Present

- Optimized 5+ production LLM-powered agents using structured prompt engineering and evaluation frameworks, improving contextual accuracy and response consistency across conversational workflows.
- Designed and deployed production-ready chatbots leveraging LLM agents as backend infrastructure, delivering scalable conversational AI solutions across multiple client implementations.
- Applied knowledge distillation techniques (QLoRA, teacher forcing) to train a lightweight in-house LLM, creating an alternative to paid external APIs for domain-specific tasks while maintaining response quality.
- Designed and evaluated LLM-driven RAG modules with vector-based retrieval (Pinecone, FAISS), delivering production-ready components within the first month.
- Architected and deployed end-to-end Shopify integration for voice AI agents, enabling real-time inventory access, order placement, and automated stock synchronization for retail customers in production environments.

miniOrange — Business Analyst Intern

Jan 2024 - May 2024

- Engineered automated data pipelines in Python and SQL with CI/CD integration, processing 1,000+ customer behavior records daily and automating payment data mapping workflows, reducing manual effort by 50-60% across teams.
- Built Tableau dashboards to visualize customer funnel metrics (MQLs, SQLs, payment conversions), enabling data-driven decision-making for product and sales teams.
- Redefined operational workflows for USA and Europe customers by engaging with 10+ clients, resulting in 5+ project scope optimizations and saving over 2 weeks of development time.

ACADEMIC PROJECTS

Text-to-SQL: Natural Language to SQL Query Generation

- Engineered and fine-tuned Flan-T5 transformer for cross-domain Text-to-SQL generation on Spider dataset (10,181 samples, 138 domains), achieving 77% valid SQL on held-out schemas and 89.8% valid SQL rate across 11 databases through schema-aware prompt engineering.
- Built and deployed production Text-to-SQL system using Streamlit and FastAPI, enabling natural language querying across arbitrary SQLite databases with automatic schema extraction and 1.6-5.5s generation latency.

Automated LLM Reasoning Evaluation Framework

- Evaluated open- and closed-source LLMs (Llama-3-8B-Instruct, Gemini-Pro 1.5) on logical reasoning benchmarks (LogiQA, ARCT) and designed a structured error taxonomy to classify reasoning-chain failures across 5 high-level and 9 fine-grained categories.
- Built an automated LLM-based evaluator achieving 81.98% accuracy, improving reasoning error detection efficiency by 30% over baseline heuristics.

AI-Med: Medical Chatbot using RAG + Vector Database

- Built and deployed a production-grade RAG pipeline using dense vector retrieval (Pinecone) to index 4,000+ pages of medical documentation, achieving 0.93 retrieval accuracy on domain-specific queries.
- Integrated LangChain with a locally hosted Meta LLaMA2 model and deployed a production-ready Flask application fusing vector-based document retrieval with dynamic LLM-based generation.

Vehicle Insurance Prediction System (Production-Grade MLOps Pipeline)

- Built a fully automated MLOps pipeline (ingestion → validation → training → deployment) with data and concept drift detection, achieving 93% accuracy (F1: 0.91, AUC: 0.92) on insurance claim predictions.
- Deployed models via AWS S3 model registry + GitHub Actions CI/CD to EC2 using Docker, Amazon ECR, and Kubernetes for scalable orchestration.
- Built secure, config-driven ML workflows with real-time logging and monitoring via Prometheus and Grafana, enabling proactive drift detection across 5 of 20 critical features.

EDUCATION

Master of Science in Computer Science

Arizona State University, United States

Aug 2024 - May 2026

CGPA: 3.96/4.0

Bachelor of Technology in Information Technology, Specialization in Data Science

Symbiosis Institute of Technology, Pune, India

May 2020 - May 2024

CGPA: 7.89/10